

# COEN 6313 Assignment 2

1<sup>st</sup> Nishant Kumar Barua  
40267821  
barua.nishant97@gmail.com

**Abstract**—This assignment implements a microservices-based architecture for a user and order management system, utilizing MongoDB for data storage and Apache Kafka for event-driven synchronization. It consists of two main microservices: the User Microservice, which handles user account creation and updates, and the Order Microservice, which manages order creation and status updates. When a user's email or delivery address is updated, an event is triggered through Kafka to synchronize the user and order databases. The API gateway routes requests to different versions of the user microservice (v1 and v2) using the Strangler Pattern, enabling gradual migration between versions. The system is deployed using Docker on a cloud VM, with configurable endpoints for API access. Optionally, GitHub integration is used for Continuous Integration and Continuous Deployment (CI/CD), automating the deployment process on code commits. This solution demonstrates scalable microservices, event-driven architecture with Kafka, and modern cloud deployment practices.

**Index Terms**—Microservice, Event-Driven Architecture, MongoDB, Message Queue, Apache Kafka, Java, SpringBoot, Strangler Pattern.

## I. IMPLEMENTATION DETAILS

We will discuss the tasks under this section, with relevant diagrams and explanations.

### A. Task 1: Creating Data Models

We have used MongoDB cloud to store our databases for the user and order data collections.

**User Data Model** We have the user data collection with fields:

- id
- firstname
- lastname
- emails
- delivery address

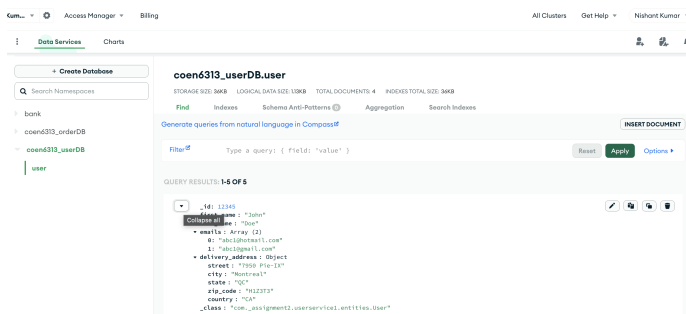


Fig. 1. User Data Model.

**Order Data Model** We have the order data collection. It has the unique orderId field. It takes the userEmails (list object) and the deliveryAddress object are of same format as the user Data model and will be used to later synchronise across the two databases. The complete list of fields are as follows:

- id
- orderId
- userEmails
- DeliveryAddress
- orderStatus
- itemList

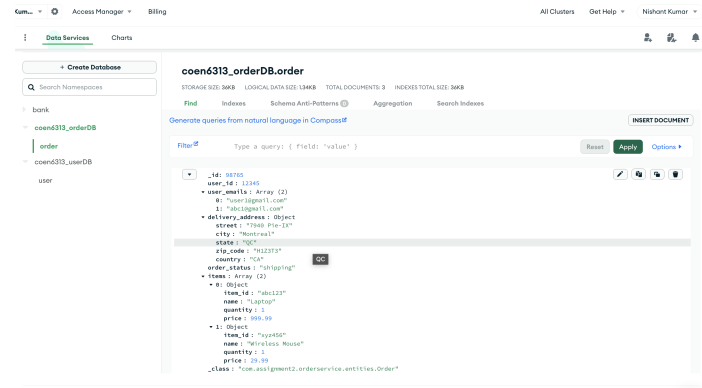


Fig. 2. Order Data Model.

### B. Task 2: Creating The Microservices

we have designed four microservices here. For the development, we have used SpringBoot to expose REST APIs for the instructed operations as detailed below.

- **User Service V1:** This microservice interacts with the User Data Model in MongoDB, using the User Repository class, which extends the Mongo Repository for this. It has endpoints to get user based on id (GET), create user (POST), update address (PUT) and update email (PUT), as exposed in the User Controller class. It operates on Port 8080. The source code is folder userserviceV1.
- **Order Service:** The order microservice interacts with the MongoDB order data base collection. It has endpoints to fetch orders by shipping status (GET), create new order (POST), update order status (PUT), delivery address (PUT) and email address (PUT). The update email address and delivery address services will be leveraged by the messaging queue to synchronise between the user and

order databases. It operates on Port 8081. The source code is folder orderservice.

- **API Gateway:** This is a SpringBoot cloud gateway which routes HTTP requests to user and order microservices. It operates on Port 80 and is the only API endpoint exposed externally for interaction with client requests. The microservice endpoints are not directly exposed for client interaction and can only be accessed by the cloud gateway. The source code is folder apigateway.

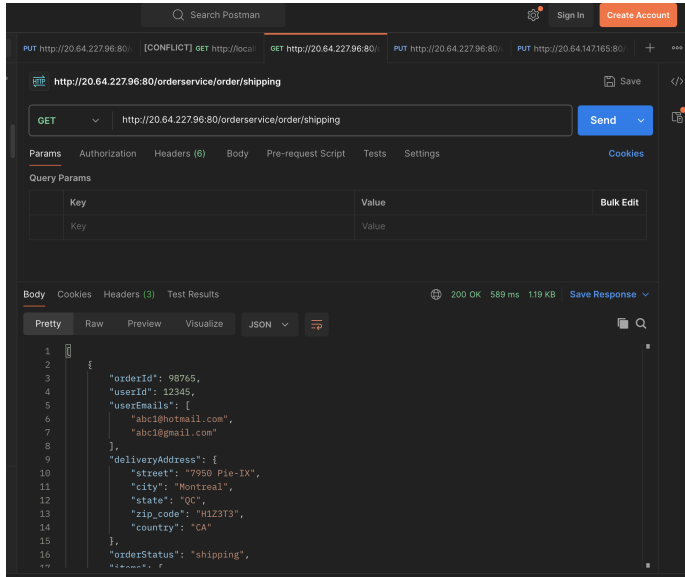


Fig. 3. Get Order EndPoint.

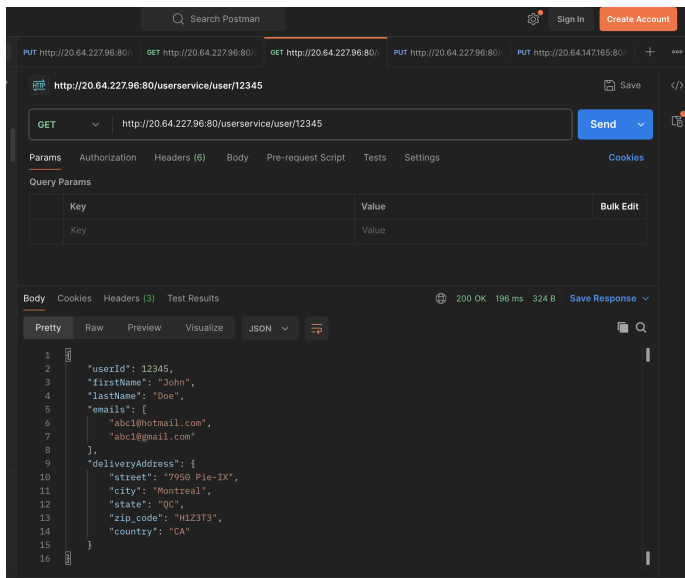


Fig. 4. Get User EndPoint.

### C. Task 3: Event Driven Architecture

We have leveraged Apache Kafka for this implementation. The Apache Kafka setup consists of a Kafka broker (Port

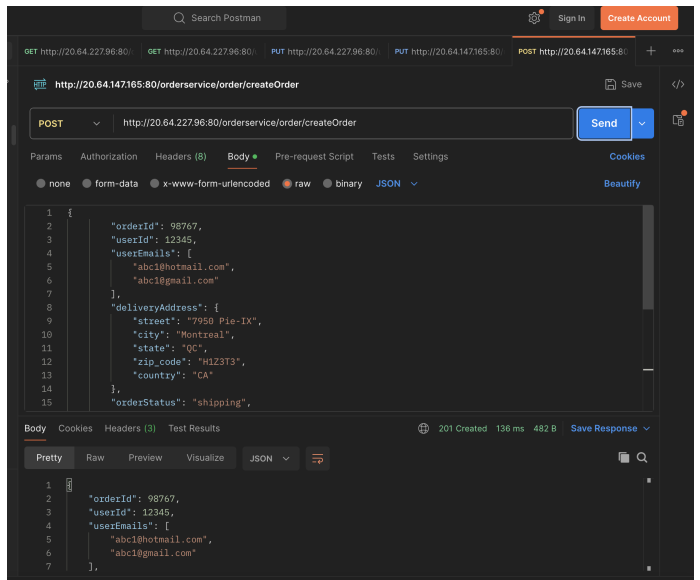


Fig. 5. Create Order EndPoint.

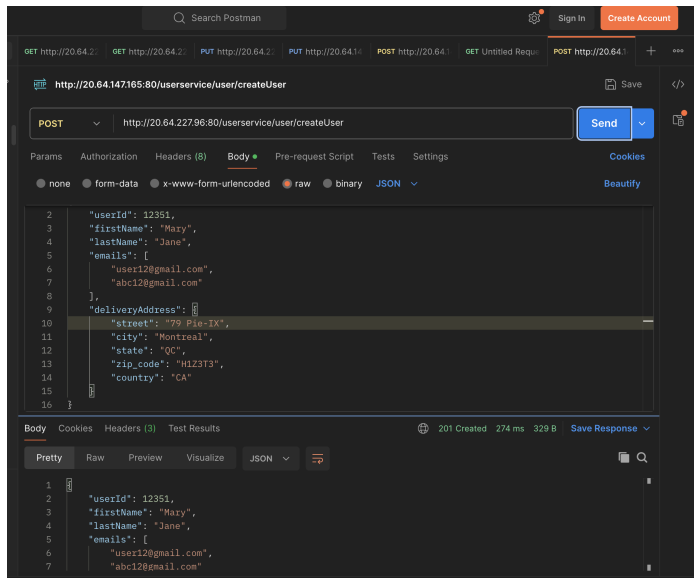


Fig. 6. Create User EndPoint.

9092) and a Zookeeper instance (Port 2181) for cluster coordination. Two Kafka topics (queues) are created: one for address updates and another for email updates. The steps are detailed below:

- The Producer microservice (both UserService V1 and V2) publish events to the respective topics when user details (address or emails) are updated via PUT requests.
- The Consumer microservice (OrderService) listen to these topics, processing the updates to synchronize the order database with the user database.

Kafka ensures reliable message delivery and event-driven communication between services, enabling decoupled data synchronization across systems.



```
Downloads — azureuser@coen6313-assignment2: ~/app — ssh -i COEN-6313-Assignment2_key.p...
=> writing image sha256:57564a65460f2d9375e48274973cb15ef232ab192964a 0.0s
=> naming to docker.io/library/app-orderservice 0.0s
=> [userservice-v1] exporting to image 1.2s
=> exporting layers 1.1s
=> writing image sha256:ada9132a23ffcd00fd32c12dacb83e1e13ffe4f88619 0.0s
=> naming to docker.io/library/app-userservice-v1 0.0s
=> [userservice-v2] exporting to image 1.1s
=> exporting layers 1.0s
=> writing image sha256:1bf051f4430f00db2684e95ed3ca0e7b212a896c9e96 0.0s
=> naming to docker.io/library/app-userservice-v2 0.0s
=> [apigateway internal] load build definition from Dockerfile 0.1s
=> transferring dockerfile: 374B 0.0s
=> [apigateway internal] load .dockerignore 0.1s
=> transferring context: 2B 0.0s
=> [apigateway internal] load build context 0.4s
=> transferring context: 34.61MB 0.3s
=> [apigateway 3/3] COPY target/*.jar app.jar 1.1s
=> [apigateway] exporting to image 0.9s
=> exporting layers 0.8s
=> writing image sha256:b38662e450479eb696f56311a7f2f0448ee60322254f4 0.0s
=> naming to assignment2.azurecr.io/apigateway:latest 0.0s
[*] Running 2/2
✓ Network app_spring-microservices Created 0.1s
✓ Container zookeeper Started 0.2s
✓ Container userservice-v1 Started 0.1s
✓ Container userservice-v2 Started 0.1s
✓ Container orderservice Started 0.2s
✓ Container apigateway Started 0.1s
✓ Container kafka Started 0.1s
azureuser@coen6313-assignment2:~/app$
```

Fig. 11. Azure Cloud VM Deployment.

### I. Task 9: CI/CD Deployment

We have used GitHub Actions to automate the CI/CD deployment on the cloud. Every time, an update is pushed to the main branch, it triggers the CI/CD pipeline to rebuild the jars and deploy the fresh instance of the app. Thus achieving end to end automation of the CI/CD pipeline. We have followed the following steps:

- First step is to build the Jars for each microservice and api gateway using maven build context.
- The, ssh into the Azure VM, build the docker containers with the Jar files.
- Final step is to deploy the containers, including Apache Kafka and Zookeeper, create the messaging queues (via Docker compose file) and run the containers.

Name (SID) and Signature	Task List		Contribution Role and Percentage ( X % )
			(If Y for the task list, write the contribution role and percentage counted for the completeness of this task)
40267821	Task 1	Y	100%
	Task 2	Y	100%
	Task 3	Y	100%
	Task 4	Y	100%
	Task 5	Y	100%
	Task 6	Y	100%
	Task 7	Y	100%
	Task 8	Y	100%
	Task 9	Y	100%

Fig. 13. Contribution.

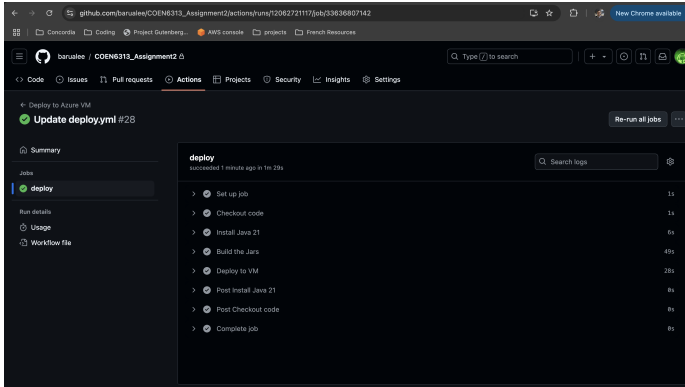


Fig. 12. CI/CD Pipeline via GitHub Actions.

## REFERENCES

- [1] Azure Documentation.
- [2] MongoDB Documentation.
- [3] SpringBoot documentation.
- [4] Docker Documentation.
- [5] GitHub Actions Documentation.